

FastVPINNs: An efficient tensor-based Python library for solving partial differential equations using hp-Variational Physics Informed Neural Networks

Thivin Anandh¹, Divij Ghose¹, and Sashikumaar Ganesan¹

¹ Department of Computational and Data Sciences, Indian Institute of Science, Bangalore, India
Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a

Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

Introduction

Partial differential equations (PDEs) are essential in modeling physical phenomena such as heat transfer, electromagnetics and fluid dynamics. The current progress in the field of scientific machine learning (SciML) has resulted in incorporating deep learning and data-driven methods to solve PDEs. Physics informed neural networks introduced by (Raissi et al., 2019) works by incorporating the governing equations of the physical systems into the Neural Network's loss function. Let us consider a two-dimensional Poisson equation as an example which reads as follows:

$$-\nabla^2 u(x) = f(x), \quad \Omega \in \mathbb{R}^2, \quad (1)$$

$$u(x) = g(x), \quad x \in \partial\Omega. \quad (2)$$

Here, $x \in \Omega$ is the spatial co-ordinate, $u(x)$ is the solution of the PDE, $f(x)$ is a known source term and $g(x)$ is the value of the solution on the domain boundary, $\partial\Omega$.

Physics-informed neural networks (PINNs), introduced by (Raissi et al., 2019), work by incorporating the governing equations of the physical systems and the boundary conditions as physics-based constraints to train the neural networks. The core idea of PINNs lies in the ability of obtaining the gradients of the solutions with respect to the input using the automatic differentiation routines available within deep learning frameworks such as tensorflow (Abadi et al., 2015). The loss function for the PINNs can be written as follows:

$$L_p(W, b) = \frac{1}{N_T} \sum_{t=1}^{N_T} |(-\Delta u_{\text{NN}}(x_t; W, b) - f(x_t))|^2,$$

$$L_b(W, b) = \frac{1}{N_D} \sum_{d=1}^{N_D} |u_{\text{NN}}(x_d; W, b) - g(x_d)|^2,$$

$$L_{\text{PINN}}(W, b) = L_p + \tau L_b.$$

where $u_{\text{NN}}(x; W, b)$ is the output of the neural network with weights W and biases b . In addition, N_T is the total number of training points in the interior of the domain Ω , N_D is the total number of training points on the boundary $\partial\Omega$, τ is a scaling factor applied to control the penalty on the boundary term and $L_{\text{PINN}}(W, b)$ is the loss function of the PINNs.

Variational Physics informed neural networks (VPINNs) (Kharazmi et al., 2019) are an extension of PINNs, where the weak form of the PDE is used to train the neural network. The weak

form of the PDE is obtained by multiplying the PDE with a test function and integrating over the domain. The method of hp-VPINNs (Kharazmi et al., 2021) was subsequently developed to increase the accuracy using h-refinement (increasing number of elements) and p-refinement (increasing the number of test functions). The loss function of hp-VPINNs with N_{elem} elements in the domain can be written as follows

$$L_v(W, b) = \frac{1}{N_{\text{elem}}} \sum_{k=1}^{N_{\text{elem}}} \left| \int_{K_k} (\nabla u_{\text{NN}}(x; W, b) \cdot \nabla v_k - f v_k) dK \right|^2,$$

$$L_b(W, b) = \frac{1}{N_D} \sum_{d=1}^{N_D} |u_{\text{NN}}(x; W, b) - g(x)|^2,$$

$$L_{\text{VPINN}}(W, b) = L_v + \tau L_b.$$

Where K_k is the k^{th} element in the domain, v_k is the test function in the corresponding element. $L_v(W, b)$ is the loss function for the weak form of the PDE and $L_{\text{VPINN}}(W, b)$ is the loss function of the hp-VPINNs. For more information on the derivation of the weak form of the PDE and the loss function of hp-VPINNs, refer to (Anandh et al., 2024).

Statement of need

The existing implementation of hp-VPINNs framework (Kharazmi, 2023) suffers from two major challenges. One being the inability of the framework to handle complex geometries and the other being the increased training time associated with the increase in number of elements within the domain. In the work (Anandh et al., 2024), we presented FastVPINNs, which addresses both of these challenges. FastVPINNs handles complex geometries by using bilinear transformation, and it uses a tensor-based loss computation to reduce the dependency of training time on number of elements. The current implementation of FastVPINNs can achieve a speed-up of upto a 100 times when compared with the existing implementation of hp-VPINNs. We have also shown that with proper hyperparameter selection, FastVPINNs can outperform PINNs both in terms of accuracy and training time, especially for problems with high frequency solutions.

Our FastVPINNs framework is built using TensorFlow 2 (Abadi et al., 2015), and provides an elegant API for users to solve both forward and inverse problems for PDEs like the Poisson, Helmholtz and Convection-Diffusion equations. The framework is well-documented with examples, which can enable users to get started with the framework with ease. The framework also provides an API to all the modules for users to write their custom functionalities. With the current level of API abstraction, users should be able to solve PDEs with less than 5 API calls as shown in the minimal working example section.

The ability of the framework to allow users to train an hp-VPINNs to solve a PDE both faster and with minimal code, can result in widespread application of this method on several real-world problems, which often require complex geometries with a large number of elements within the domain.

Modules

The FastVPINNs framework consists of five core modules, which are:

- Geometry - This module handles mesh generation.
- FE - This module handles the finite element test functions and their gradients.
- Data - This module handles tensor assembly.

- Physics - This module handles the tensor-based loss computation for the weak form of the PDE.
- Model - This module handles the dense neural network for training the PDE.

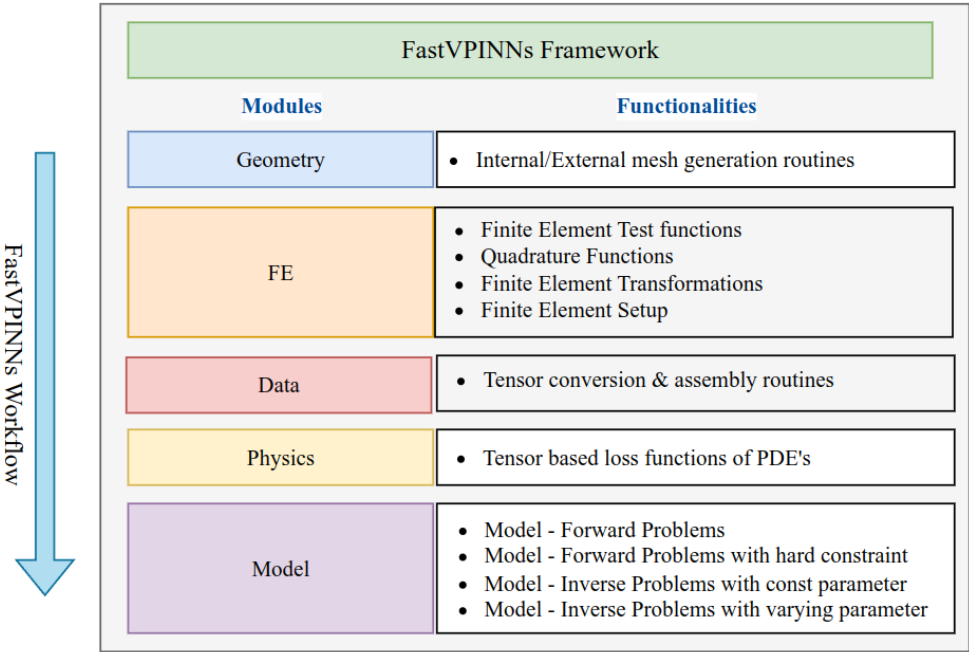


Figure 1: FastVPINNs Modules

Geometry Module

This module provides the functionality to define the geometry of the domain. The user can either generate a quadrilateral mesh internally or read an external .mesh file. The module also provides the functionality to obtain boundary points for complex geometries.

FE Module

The FE module is responsible for handling the finite element test functions and their gradients. The module's functionality can be broadly classified into into four categories:

- Finite element test functions:** Provides the test function values and its gradients for a given reference element.
- Quadrature functions:** Provides the quadrature points and weights for a given element based on the quadrature order selected by the user. These values will be used for numerical integration of the weak form of the PDE.
- Transformation functions:** Provides the implementation of transformation routines such as 'Affine' transformation and 'Bilinear' transformation. This can be used to transform the test function values and gradients from the reference element to the actual element.
- Finite Element Setup:** Provides the functionality to set up the test functions, quadrature rules and the transformation for every element and save them in a common wrapper to access them. Further, it also hosts functions to plot the mesh with quadrature points, assign boundary values based on the boundary points obtained from the geometry module, calculation of the forcing term etc.

89 **Data Module:**

90 The Data module collects data from all modules which are required for training and converts
91 them to a tensor data type with required precision. It also assembles the test function values and
92 gradients to form a three-dimensional tensor, which will be used during the loss computation.

93 **Physics Module:**

94 This module contains functions which are used to compute the variational loss function for
95 different PDEs. These functions accept the test function tensors and the predicted gradients
96 from the neural network along with the forcing matrix to compute the PDE residuals.

97 **Model Module:**

98 This module contains custom subclasses of the `tensorflow.keras.Model` class, which are
99 used to train the neural network. The module provides the functionality to train the neural
100 network using the tensor based loss computation and also provides the functionality to predict
101 the solution of the PDE for a given input.

102 **Minimal Working Example**

103 With the higher level of abstraction provided by the FastVPINNs framework, users can solve
104 a PDE with just 5 API calls. Shown below is a minimal working example to solve a Poisson
105 equation using the FastVPINNs framework.

```
#load the geometry
domain = Geometry_2D("quadrilateral", "internal",
                     100, 100, ".")
cells, boundary_points =
    domain.generate_quad_mesh_internal(\
        x_limits=[0, 1], y_limits=[0, 1],\
        n_cells_x=4, n_cells_y=4,\
        num_boundary_points=400)

# Note: Users need to provide the necessary
#       boundary values and the forcing function

# load the FESpace
fespace = Fespace2D(domain.mesh, cells, boundary_points, \
                    domain.mesh_type, fe_order=5, \
                    fe_type="jacobi", quad_order=5, \
                    quad_type="legendre", \
                    fe_transformation_type="bilinear" \
                    bound_function_dict=bound_function_dict, \
                    bound_condition_dict=bound_condition_dict, \
                    forcing_function=rhs, \
                    output_path=i_output_path, \
                    generate_mesh_plot=True)

# Instantiate Data handler
dh = datahandler2D(fespace, domain, dtype=tf.float32)

# Instantiate the model with the loss function for the model
model = DenseModel(layer_dims=[2, 30, 30, 30, 1], \
                   learning_rate_dict=0.01, params_dict=params_dict,
```

```

## Loss function of poisson2D
loss_function=pde_loss_poisson,
input_tensors_list=\
    [in_tensor,
     dir_in,
     dir_out],
orig_factor_matrices=\
    [dh.shape_val_mat_list,
     dh.grad_x_mat_list,
     dh.grad_y_mat_list],
force_function_list=dh.forcing_function_list,\
tensor_dtype=tf.float32,
use_attention=i_use_attention,
activation=i_activation,
hessian=False)

# Train the model
for epoch in range(1000):
    model.train_step()

```

Testing

The FastVPINNs framework has a strong testing suite, which tests the core functionalities of the framework. The testing suite is built using the pytest framework and is integrated with the continuous integration pipeline provided by Github Actions. The testing functionalities can be broadly classified into the three categories:

- **Unit Testing:** Covers the testing of individual modules and their functionalities.
- **Integration Testing:** Covers the overall flow of the framework. Different PDEs are solved with different parameters to check if the accuracy after training is within the acceptable limits. This ensures that the collection of modules work together as expected.
- **Compatibility Testing:** The framework is tested with different versions of python such as 3.8, 3.9, 3.10, 3.11 and on different versions of OS such as Ubuntu, MacOS and Windows to ensure the compatibility of the framework across different platforms.

Acknowledgements

We thank Shell Research, India, for providing partial funding for this project. We are thankful to the MHRD Grant No.STARS-1/388 (SPADE) for partial support.

References

- Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., Corrado, G. S., Davis, A., Dean, J., Devin, M., Ghemawat, S., Goodfellow, I., Harp, A., Irving, G., Isard, M., Jia, Y., Jozefowicz, R., Kaiser, L., Kudlur, M., ... Zheng, X. (2015). *TensorFlow: Large-scale machine learning on heterogeneous systems*. <https://www.tensorflow.org/>
- Anandh, T., Ghose, D., Jain, H., & Ganesan, S. (2024). *FastVPINNs: Tensor-driven acceleration of VPINNs for complex geometries*. <https://arxiv.org/abs/2404.12063>
- Kharazmi, E. (2023). *hp-VPINNs: High-performance variational physics-informed neural networks*. <https://github.com/ehsankharazmi/hp-VPINNs>

- 130 Kharazmi, E., Zhang, Z., & Karniadakis, G. E. (2019). Variational physics-informed neural
131 networks for solving partial differential equations. *arXiv Preprint arXiv:1912.00873*.
- 132 Kharazmi, E., Zhang, Z., & Karniadakis, G. E. (2021). hp-VPINNs: Variational physics-
133 informed neural networks with domain decomposition. *Computer Methods in Applied
134 Mechanics and Engineering*, 374, 113547.
- 135 Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks:
136 A deep learning framework for solving forward and inverse problems involving nonlinear
137 partial differential equations. *Journal of Computational Physics*, 378, 686–707.

DRAFT